

The Go concurrency model

A tour on simplicity

Mario Macías, Ph.D.

New Relic Inc

Assembler School

February 2020

mmacias@newrelic.com

[@MaciasUPC](https://twitter.com/MaciasUPC)

<http://macias.info>

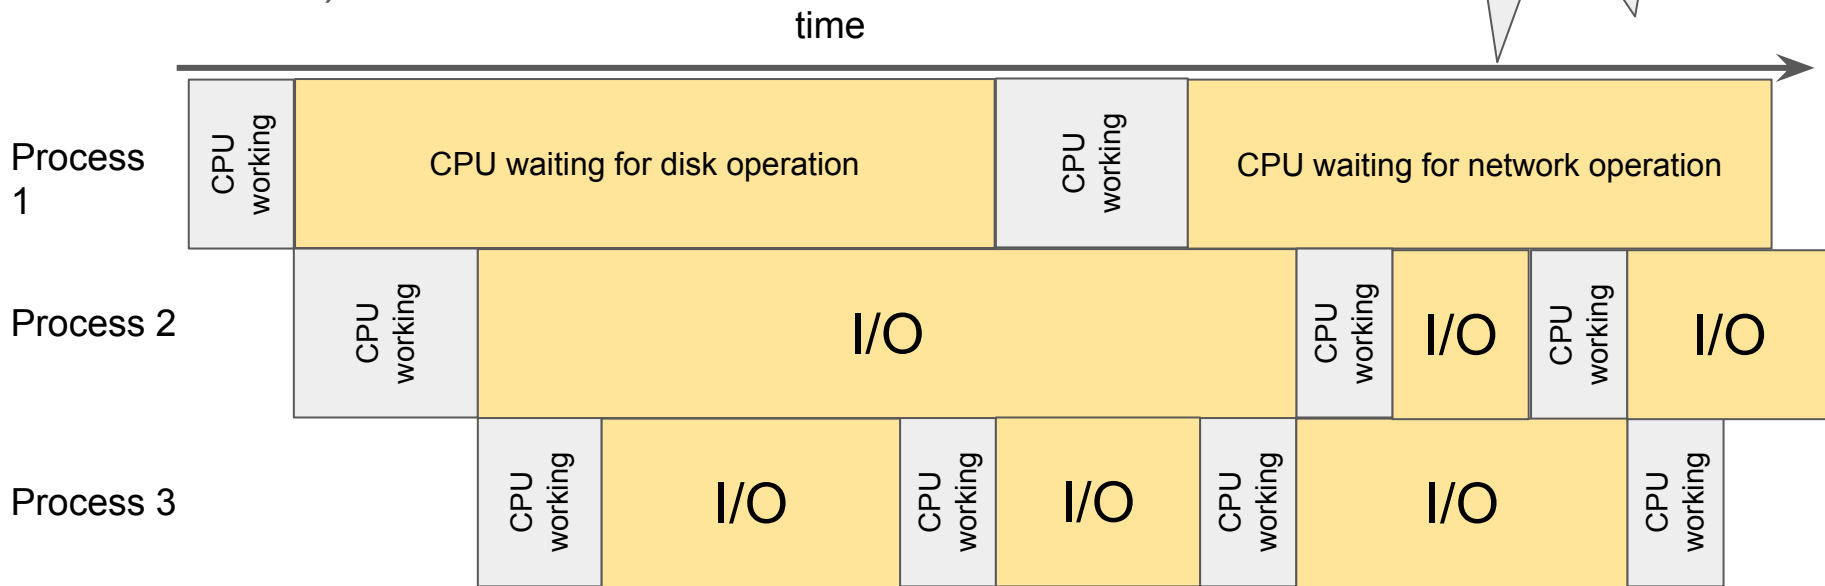
Who am I?



Why multiprocessing?


- CPU becomes idle on I/O operations (e.g. accessing the disk)

yey! You could use this idle time to work on other things!!

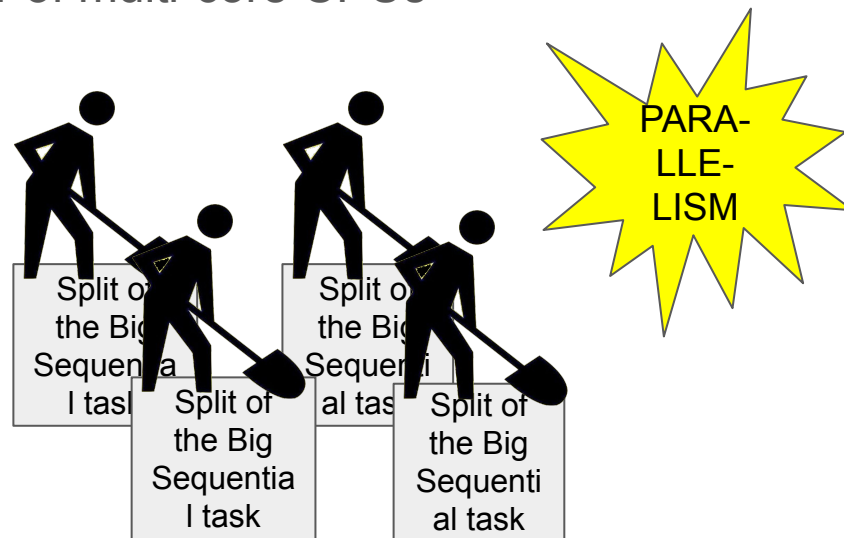


Why multiprocessing?

- Battle for GigaHertz ended
- Programs need to squeeze the power of multi-core CPUs



single-core times



Multicore approach

Terminology

- Processes

- A program in execution: OS can spawn multiple processes in parallel
- It has its own isolated space in the memory
- Requires heavyweight management from the OS

- Threads

- A process can spawn multiple threads
- They share the process memory, but the stack
- Lighter management than Processes

- Parallelism

- Dividing a usually big task in multiple subtasks (e.g. threads or processes) that are executed at the same time (e.g. different cores)

- Concurrency

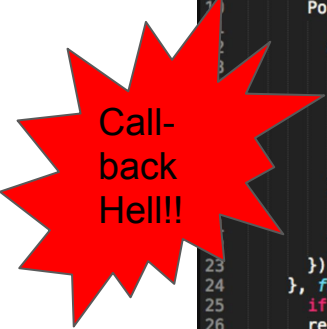
- The problem of effectively sharing resources and data between parallel tasks
- The problem of effectively coordinating parallel tasks

Threads limitations

- Despite they are lighter than processes, they are still heavyweight
 - *time for demo!*
-

Lightweight concurrency: async, callback-based

- Require a "runtime" that schedules the tasks. It can be a library.
- Requires stop coding as most languages are designed for



Call-back Hell!!!

```
1
2 var async = require("async");
3
4 User.find(userId, function(err, user){
5   if (err) return errorHandler(err);
6   User.all({where: {id: {$in: user.friends}}}, function(err, friends) {
7     if (err) return errorHandler(err);
8     async.each(friends, function(friend, done){
9       friend.posts = [];
10      Post.all({where: {userId: {$in: friend.id}}}, function(err, posts) {
11        if (err) return errorHandler(err);
12        async.each(posts, function(post, donePosts){
13          friend.push(post);
14          Comments.all({where: post.id}, function(err, comments) {
15            if (err) donePosts(err);
16            post.comments = comments;
17            donePosts();
18          });
19        }, function(err) {
20          if (err) return errorHandler(err);
21          done();
22        });
23      });
24    }, function(err) {
25      if (err) return errorHandler(err);
26      render(user, friends);
27    });
28  }
29 });
30
```

Lightweight concurrency: coroutines

- The programmer needs to state when the execution can go to another part of the code
- The runtime needs to be part of the language implementation

```
var test = coroutine(function*() {  
    console.log('Hello!');  
    var x = yield;  
    console.log('First I got: ' + x);  
    var y = yield;  
    console.log('Then I got: ' + y);  
});  
// prints 'Hello!'  
test('a dog'); // prints 'First I got: a dog'  
test('a cat'); // prints 'Then I got: a cat'
```


Goroutines

- Simplifies even more the approach of coroutines
- The programmer doesn't need to put "yields" in the code → Cleaner code
- The Go API and implementation is full of "hidden" yields that gives the control to the Go runtime where it's more efficient to do it
- Drawback: you pay the overhead in performance even if you don't use them

Threads vs Fibers

- Threads and processes: preemptive multitasking
 - The "steals" the CPU to the processes/threads when they perform given operations, or after a given time.
 - If a process/thread hogs the CPU, the OS steals it and schedules another process.
 - It has a big cost: the OS needs to save the complete CPU status for the paused process and restore the status of the new allocated process.
- Fibers (a.k.a. lightweight threads): cooperative multitasking
 - An async callback function or a coroutine are types of fibers
 - The "runtime" allocates a coroutine/function to be executed in one of the threads that run the runtime
 - The coroutines return voluntarily the execution → they only need to save the important data for the later use
 - If a coroutine hogs the CPU and "doesn't cooperate", the other coroutines aren't executed

Launching a Goroutine

```
go nameOfTheFunction(arguments)
```

or

```
go func() {  
    /* your instructions here */  
}()
```

Need of syncing goroutines

- The above goroutines won't usually complete (demo time)

```
for i := 0; i < 10 ; i++ {  
    numTask := i  
    go func() {  
        fmt.Printf("Running parallel task %d\n", numTask)  
    }()  
}  
fmt.Println("All the parallel tasks have ended. Exiting now")
```

WaitGroup

- WaitGroup allows blocking a goroutine until a group of other goroutines finish
- `wg := sync.WaitGroup{}`
`wg.Add(<num>)`
 - Instantiates a WaitGroup
 - Sets the num of parallel goroutines we will need to wait for
- `wg.Done()`
 - Invoked from each goroutine
 - Decreases the waitgroup counter, meaning that this goroutine has finished
- `wg.Wait()`
 - Usually invoked from the instantiator of the WaitGroup
 - The execution is blocked until all the <num> goroutines invoke `wg.Done()`

Mutual Exclusion

- Demo time

Mutual Exclusion

- Sometimes we need to make sure that a given segment of the code is accessed only by a goroutine at the same time
- `m := sync.Mutex{}`
 - Instantiates a mutual exclusion
- `m.Lock()`
 - Gets the exclusive permission to access the code after this invocation
 - If another goroutine already has this permission, it waits
- `m.Unlock()`
 - Releases the permission of the Mutex so other goroutines can get it
- Alternative implementation
 - `sync.RWMutex{}`
 - Multiple Readers can read at the same time.
 - One single Writer locks all the accesses

Other "classic" tools

- Package "atomic"
 - Allows some atomic operations on basic types (Read and Set)
- `sync.Once{}`
 - Makes sure that a function is only executed once, even if invoked from multiple goroutines
- `sync.Pool{}`
 - Allows reusing a set of items that may be expensive to instantiate
 - E.g. Database Connections

Channels

- The most used synchronization tool from Go
- It allows sending and receiving values from different goroutines
- `ch := make(chan string)`
 - Creates a channel that shares strings
- `ch <- "hello"`
 - Sends a string to a channel
- `str := <-ch`
 - Receives a string via a channel
- `close(ch)`
 - Closes the channel

Unbuffered channel

- The receiver of a channel always block until there is something in the channel
- The sender of a channel is also blocked until the item is received
- Warn!
 - Potential deadlocks

Buffered channel

```
ch := make(chan string, 100)    // buffer length: 100 strings
```

- The receiver via a channel will block until something can be received
 - When the element is received, is removed from the buffer
- The sender:
 - If the buffer is not full, it will send the value and continue its execution
 - If the buffer is full, it will block until another element is removed

Close-based synchronization

- When a channel is closed, the receiver won't block anymore and will continue its execution, without having received anything

Multiple receivers

- Multiple goroutines can send/receive at the same time
- There is no guarantee of a fair use

Channel iteration

- A channel can be used in a "range" clause inside a loop
- It is a comfortable way to say "keep receiving values until it is closed"

```
// sweets is a channel
```

```
for sweet := range sweets {  
    fmt.Println("Received a delicious", sweet)  
}
```

Channel select

```
select {  
  case v := <- foo:  
    fmt.Println("received from foo: ", v)  
  case v := <- bar:  
    fmt.Println("from bar:", v)  
}
```

- The following code will block until one of the channels has data
- I will receive data from only one channel

channel timeouts

- `time.After(< duration >)` returns a channel that sends a value after the duration
- With the previous "select" clause, it can implement a channel timeout

```
select {  
case v := <- foo:  
    fmt.Println("received from foo: ", v)  
case <- time.After(5 * time.Second):  
    fmt.Println("timeout while waiting for a foo")  
}
```


The Go concurrency model

Thank you for your attention!

Mario Macías, Ph.D.

New Relic Inc

Assembler School

February 2020

mmacias@newrelic.com

[@MaciasUPC](https://twitter.com/MaciasUPC)

<http://macias.info>